

CSCI 570 - Spring 2022 - HW 4

Due: Feb 9th

Section 1: Heaps

Reading Assignment: Kleinberg and Tardos, Chapter 2.5.

Problem 1

[10 points] Design a data structure that has the following properties (assume n elements in the data structure, and that the data structure properties need to be preserved at the end of each operation):

- Find median takes $\mathcal{O}(1)$ time
- Insert takes $\mathcal{O}(\log n)$ time

Do the following:

1. Describe how your data structure will work.
2. Give algorithms that implement the Find-Median() and Insert() functions.

Solution. We use the $dn/2e$ smallest elements to build a max-heap and use the remaining $bn/2c$ elements to build a min-heap. The median will be at the root of the max-heap and hence accessible in time $\mathcal{O}(1)$ (we assume the case of even n , median is $n/2$ -th element when elements are sorted in increasing order).

Insert() algorithm: For a new element x

- Initialize len of maxheap (maxlen) and len of minheap (minlen) to 0.
- Compare x to the current root of max-heap.
- If $x < \text{median}$, we insert x into the max-heap. Maintain length of maxheap say maxlen, and every time you insert element into maxheap increase maxlen by 1. Otherwise, we insert x into the min-heap, and increase length of minheap say minlen by 1. This takes $\mathcal{O}(\log n)$ time in the worst case.
- If $\text{size}(\text{maxHeap}) > \text{size}(\text{minHeap}) + 1$, then we call Extract-Max() on max-heap, and decrease maxlen by 1 of and insert the extracted value into the min-heap and increase minlen by 1. This takes $\mathcal{O}(\log n)$ time in the worst case.

- Also, if $\text{size}(\text{minHeap}) > \text{size}(\text{maxHeap})$, we call $\text{Extract-Min}()$ on min-heap, decrease minlen by 1 and insert the extracted value into the max-heap and increase maxlen by 1. This takes $\mathcal{O}(\log n)$ time in the worst case.

Find-Median() algorithm:

- If $(\text{maxlen} + \text{minlen})$ is even: return $(\text{sum of roots of max heap and min heap})/2$ as median
- Else if $\text{maxlen} > \text{minlen}$: return root of max heap as median
- Else: return root of min heap as median



Rubric. • 2 pt: Using max heap and min heap to store first half and second half of elements.

- 3 pt: Comparing element to the root and proper if conditions for inserting in appropriate heap.
- 2 pt: Proper if conditions for returning the median.
- 3 pt: Correct Time Complexity



Section 2: MST

Reading Assignment: Kleinberg and Tardos, Chapter 4.5.

Problem 2

[10 points] Let us say that a graph $G = (V, E)$ is a near tree if it is connected and has at most $n+8$ edges, where $n = |V|$. Give an algorithm with running time $\mathcal{O}(n)$ that takes a near tree G with costs on its edges, and returns a minimum spanning tree of G . You may assume that all edge costs are distinct.

- Solution.**
1. To do this, we apply the Cycle Property nine times. That is, we perform BFS until we find a cycle in the graph G , and then we delete the heaviest edge on this cycle.
 2. We have now reduced the number of edges in G by one while keeping G connected and (by the Cycle Property) not changing the identity of the minimum spanning tree.
 3. If we do this a total of nine times, we will have a connected graph H with $n - 1$ edges and the same minimum spanning tree as G . But H is a tree, and so in fact it is the minimum spanning tree.

4. The running time of each iteration is $\mathcal{O}(m+n)$ for the BFS and subsequent check of the cycle to find the heaviest edge; here $m \leq n + 8$, so this is $\mathcal{O}(n)$.



Rubric. • 3 pt: Applying cycle property 9 times

- 2 pt: Figuring out finding LCA to find heaviest edge
- 2 pt: Figuring out that we will have to delete the heaviest edge on this cycle
- 3 pt: Correct Time Complexity



Problem 3

[12 points] Considering the following graph G :




Figure 1: Graph G .

1. [4 points] In graph G , if we use Kruskal's Algorithm to find the MST, what is the third edge added to the solution? Select all correct answers
 - a. E-F
 - b. D-E
 - c. A-B
 - d. C-F
 - e. D-F

2. [4 points] In graph G , if we use Prim's Algorithm to find MST starting at A , what is the second edge added to the solution?
- a. B-G
 - b. B-E
 - c. D-E
 - d. A-D
 - e. E-F
3. [4 points] What is the cost of the MST in the Graph?
- a. 18
 - b. 19
 - c. 20
 - d. 21
 - e. 22

Solution. 1. c. A-B

2. b. B-E

3. c. 20



Rubric. 4 points for each correct selection



Problem 4

[10 points] A network of n servers under your supervision is modeled as an undirected graph $G = (V, E)$ where a vertex in the graph corresponds to a server in the network and an edge models a link between the two servers corresponding to its incident vertices. Assume G is connected. Each edge is labeled with a positive integer that represents the cost of maintaining the link it models. Further, there is one server (call its corresponding vertex as S) that is not reliable and likely to fail. Due to a budget cut, you decide to remove a subset of the links while still ensuring connectivity. That is, you decide to remove a subset of E so that the remaining graph is a spanning tree. Further, to ensure that the failure of S does not affect the rest of the network, you also require that S is connected to exactly one other vertex in the remaining graph.

Design an algorithm that given G and the edge costs efficiently decides if it is possible to remove a subset of E , such that the remaining graph is a spanning tree where S is connected to exactly one other vertex and (if possible) finds a solution that minimizes the sum of maintenance costs of the remaining edges.

Solution. First we need to check the possibility of node S having only one neighbor in a spanning tree of the underlying graph. The best way to check this is to remove node S and all its adjacent edges to form a graph G' . If G' is a connected graph, then we can claim it is possible to have a spanning tree where node S has only one neighbor. To check the connectivity of G' , the simplest way is to run a DFS or BFS algorithm on G' .

Considering that G' is connected, we need to find the spanning tree that minimizes the maintenance cost. Therefore we need to find the MST with the additional constraint that S should be a leaf. Therefore, we remove S and all its adjacent edges to form G' . We run Prim's (or any other MST algorithm) to find the MST of G' . Among all edges adjacent to S , we find the one with the minimum maintenance cost and connect S to the MST using this edge. The resulting graph will still be a spanning tree and in this spanning tree S will be a leaf. ■

Rubric. • 5 pt: Building your spanning tree in a way that node S is a leaf

• 5 pt: Making sure that the final MST is connected ■

Section 3: Shortest Path

Reading Assignment: Kleinberg and Tardos, Chapter 4.4.

Problem 5

[20 points] Given a connected graph $G = (V, E)$ with positive edge weights. In V , s and t are two nodes for shortest path computation, prove or disprove with explanations:

1. If all edge weights are unique, then there is a single shortest path between any two nodes in V .
2. If each edge's weight is increased by k , the shortest path cost between s and t will increase by a multiple of k .
3. If the weight of some edge e decreases by k , then the shortest path cost between s and t will decrease by at most k .
4. If each edge's weight is replaced by its square, i.e., w to w^2 , then the shortest path between s and t will be the same as before but with different costs.

Solution. 1. False. Counter example: (s, a) with weight 1, (a, t) with weight 2 and (s, t) with weight 3. There are two shortest path from s to t though the edge weights are unique.

Rubric. (a) 2 pt: Correct T/F claim
 (b) 3 pt: Provides a correct counterexample as explanation

■

2. False. Counter example: suppose the shortest path $s \rightarrow t$ consist of two edges, each with cost 1, and there is also an edge $e = (s, t)$ in G with $\text{cost}(e)=3$. If now we increase the cost of each edge by 2, e will become the shortest path (with the total cost of 5).

Rubric. (a) 2 pt: Correct T/F claim
 (b) 3 pt: Provides a correct counterexample as explanation

■

3. False.
- Only true when we have the assumption that after decreasing all edge weights are still positive (however we don't have this in the problem). For any two nodes s, t , assume that P_1, \dots, P_k are all the paths from s to t . If e belongs to P_i then the path cost decrease by k , otherwise the path cost unchanged. Hence all paths from s to t will decrease by at most k . As shortest path is among them, then the shortest path cost will decrease by at most k .
 - When 1) there is cycle in the graph, 2) and there is a path from s to t that goes through that cycle, 3) and after decreasing, the sum of edge weights of that cycle becomes negative, then the shortest path from s to t will go over the cycle for infinite times, ending up with infinite path cost, hence not "decrease by at most k ".

Rubric. (a) 2 pt: Correct T/F claim
 (b) 3 pt: Considers the negative cycle case

■

4. False. Counter example: 1) suppose the original graph G composed of $V = \{A, B, C, D\}$ and $E : (A \rightarrow B) = 100, (A \rightarrow C) = 51, (B \rightarrow D) = 1, (C \rightarrow D) = 51$, then the shortest path from A to D is $A \rightarrow B \rightarrow D$ with length 101. 2) After squaring this path length become $100^2 + 1^2 = 10001$. However, $A \rightarrow C \rightarrow D$ has path length $51^2 + 51^2 = 5202 < 10001$ Thus $A \rightarrow C \rightarrow D$ become the new shortest path from A to D .

Rubric. (a) 2 pt: Correct T/F claim
 (b) 3 pt: Provides a correct counterexample as explanation

■

■

Problem 6

[16 points] Consider a directed, weighted graph G where all edge weights are positive. You have one Star, which allows you to change the weight of any one edge to zero. In other words, you may change the weight of any one edge to zero. Propose an efficient method based on Dijkstra's algorithm to find a lowest-cost path from node s to node t , given that you may set one edge weight to zero.

Note: you will receive 10 points if your algorithm is efficient. This means your method must do better than the naive solution, where the weight of each node is set to 0 per time and the Dijkstra's algorithm is applied every time for lowest-cost path searching. You will receive full points (16 points) if your algorithm has the same run time complexity as Dijkstra's algorithm.

Solution. Use Dijkstra's algorithm to find the shortest paths from s to all other vertices. Reverse all edges and use Dijkstra to find the shortest paths from all vertices to t . Denote the shortest path from u to v by $u \rightsquigarrow v$, and its length by $\delta(u, v)$. Now, try setting each edge to zero. For each edge $(u, v) \in E$, consider the path $s \rightsquigarrow u \rightarrow v \rightsquigarrow t$. If we set $w(u, v)$ to zero, the path length is $\delta(s, u) + \delta(v, t)$. Find the edge for which this length is minimized and set it to zero; the corresponding path $s \rightsquigarrow u \rightarrow v \rightsquigarrow t$ is the desired path. The algorithm requires two invocations of Dijkstra, and an additional $\mathcal{O}(E)$ time to iterate through the edges and find the optimal edge to take for free. Thus the total running time is the same as that of Dijkstra:

- If we implement Dijkstra's algorithm with Fibonacci heap
 - time complexity of both the Dijkstra's algorithm and the proposed method is $\mathcal{O}(E + V \log V)$.
 - the corresponding time complexity of naive is $\mathcal{O}(E(E + V \log V))$
- If we implement Dijkstra's algorithm with a binary heap
 - the complexity of Dijkstra's algorithm is $\mathcal{O}((E + V) \log V)$, so that the proposed method is $\mathcal{O}(2(E + V) \log V + E)$, since $E = \mathcal{O}(E \log V)$, then it has same time complexity as Dijkstra's algorithm, i.e., $\mathcal{O}((E + V) \log V)$.
 - the corresponding time complexity of naive is $\mathcal{O}(E(E + V) \log V)$



Rubric. 1. Time complexity analysis: all kinds of implementations for Dijkstra's algorithm should be fine.

- 6 pt for correct analysis based on one specific implementation approach for the Dijkstra's algorithm
- 0 otherwise

2. Approach description: based on the actual time complexity (the student's time complexity is not necessarily correct)
- 10 pt if the approach is of the same complexity as Dijkstra's algorithm
 - 4 pt if the approach is more efficient than naive but not has the same complexity as Dijkstra's algorithm
 - 0 pt if the naive approach or any other approach with larger complexity than naive is proposed

